

Particle Filters

16-833 Robot Localization and Mapping

Fall 2024

Montiel Abello

Slides adapted from Eric Westman

State Estimation

Parametric Methods


- Model state using a parametric distribution (e.g. normal distribution)
- Can give optimal estimate if assumptions hold
- Examples: Kalman filter, EKF, factor graph optimization

Non-parametric Methods

- Do not assume a particular model
- Allows tracking arbitrary distributions
- Often use particles or kernels to represent underlying distribution

Particle Filter Applications

- Localization (focus in this class)
- SLAM (e.g. FASTSLAM algorithm)
- State estimation, generally




The Big Idea

Use particles as **samples** of the distribution that represents our **belief** of the state


The Big Idea

Use particles as **samples** of the distribution that represents our **belief** of the state



Background...

Sampling



$$I(x) = I\left(\begin{bmatrix} u \\ v \end{bmatrix}\right) = \begin{cases} 1 & \text{if } x \in S \\ 0 & \text{if } x \notin S \end{cases}$$

Approximate Area:

- sample N $x_i \in X$
- Area $\approx \frac{1}{N} \sum_{i=1}^N I(x_i) a \cdot b$

Assumptions:


- have indicator function
- have method of drawing uniformly distributed random numbers

Why is sampling nontrivial?


- Bernoulli distribution: $x \in \{0, 1\}$
 $P(x=1) = p$
- Uniform distribution: $x \sim \text{Uni}(0, 1)$
 - hardware, pseudorandom number generators
 - allow us to sample more complex distributions

Inverse Transform Sampling Method

- e.g. sample 0-mean Gaussian



- ① find or approx CDF
$$CDF(x) = \int_{-\infty}^x P(t) dt$$




- ② generate uniformly distributed random no. $u_i \sim \text{Uni}(0,1)$
$$x_i = CDF^{-1}(u_i)$$

Problem: can't do this for multivariate distributions

Rejection Sampling

- Goal: sample target distribution $p(x)$ (can evaluate, can't sample)
- Known: proposal distribution $q(x)$ (can sample, can evaluate)
e.g. uniform, gaussian



- Set up: Choose M s.t. $p(x) \leq Mq(x)$ over entire support of $p(x)$

- ① sample $x_i \sim q(x)$ and $u_i \sim \text{Uni}(0,1)$
- ② if $u_i < \frac{p(x_i)}{Mq(x_i)}$, accept x_i

else reject x_i , go to ①

\Rightarrow densely sample where $p(x)$ close to $Mq(x)$ i.e. where probability density higher

Problem: in high dim, need $M \uparrow$
hard to get good fit

Importance Sampling

• Notation: $x \sim p(x)$

$$E[x] = \int_{-\infty}^{\infty} x p(x) dx \\ \triangleq E_{p(x)}[x]$$

$$E_{p(x)}[f(x)] = \int_{-\infty}^{\infty} f(x) p(x) dx$$

- target $p(x)$ - can't sample
- proposal $q(x)$ - can sample

$$\begin{aligned} E_{p(x)}[f(x)] &= \int f(x) p(x) dx \\ &= \int f(x) p(x) \frac{q(x)}{q(x)} dx \\ &= \int \frac{p(x)}{q(x)} f(x) q(x) dx \\ &= E_{q(x)} \left[\frac{p(x)}{q(x)} f(x) \right] \end{aligned}$$



• now generate samples $x_i \sim q(x)$

$$E_{p(x)}[f(x)] \sim \frac{1}{N} \sum_{i=1}^N \frac{p(x_i)}{q(x_i)} f(x_i)$$

• need to know ratio $\frac{p(x)}{q(x)}$



Importance Sampling

$$\begin{aligned}\mathbb{E}_{p(x)} [f(x)] &= \int_{-\infty}^{\infty} f(x) p(x) dx \\ &= \int_{-\infty}^{\infty} f(x) p(x) \frac{q(x)}{q(x)} dx \\ &= \int_{-\infty}^{\infty} \frac{p(x)}{q(x)} f(x) q(x) dx \\ &= \mathbb{E}_{q(x)} \left[\frac{p(x)}{q(x)} f(x) \right]\end{aligned}$$





Importance Sampling

$$\begin{aligned} \mathbb{E}_{p(x)} [f(x)] &= \int_{-\infty}^{\infty} f(x) p(x) dx \\ &= \int_{-\infty}^{\infty} f(x) p(x) \frac{q(x)}{q(x)} dx \\ &= \int_{-\infty}^{\infty} \frac{p(x)}{q(x)} f(x) q(x) dx \\ &= \mathbb{E}_{q(x)} \left[\frac{p(x)}{q(x)} f(x) \right] \end{aligned}$$





Importance Sampling

$$\begin{aligned} \mathbb{E}_{p(x)} [f(x)] &= \int_{-\infty}^{\infty} f(x) p(x) dx \\ &= \int_{-\infty}^{\infty} f(x) p(x) \frac{q(x)}{q(x)} dx \\ &= \int_{-\infty}^{\infty} \frac{p(x)}{q(x)} f(x) q(x) dx \\ &= \mathbb{E}_{q(x)} \left[\frac{p(x)}{q(x)} f(x) \right] \end{aligned}$$






Importance Sampling

$$\begin{aligned} \mathbb{E}_{p(x)} [f(x)] &= \int_{-\infty}^{\infty} f(x) p(x) dx \\ &= \int_{-\infty}^{\infty} f(x) p(x) \frac{q(x)}{q(x)} dx \\ &= \int_{-\infty}^{\infty} \frac{p(x)}{q(x)} f(x) q(x) dx \\ &= \mathbb{E}_{q(x)} \left[\frac{p(x)}{q(x)} f(x) \right] \end{aligned}$$



Importance Sampling


$$\begin{aligned} \mathbb{E}_{p(x)} [f(x)] &= \int_{-\infty}^{\infty} f(x) p(x) dx \\ &= \int_{-\infty}^{\infty} f(x) p(x) \frac{q(x)}{q(x)} dx \\ &= \int_{-\infty}^{\infty} \frac{p(x)}{q(x)} f(x) q(x) dx \\ &= \mathbb{E}_{q(x)} \left[\frac{p(x)}{q(x)} f(x) \right] \end{aligned}$$



Sampling Application: Neural Rendering

Inverse transform sampling:



- NeRF: Representing scenes as neural radiance fields for view synthesis [Mildenhall et al., 2021]
- Mip-NeRF: A Multiscale Representation for Anti-Aliasing Neural Radiance Fields [Barron et al., 2021]



Sampling Application: Neural Rendering

Inverse transform sampling:

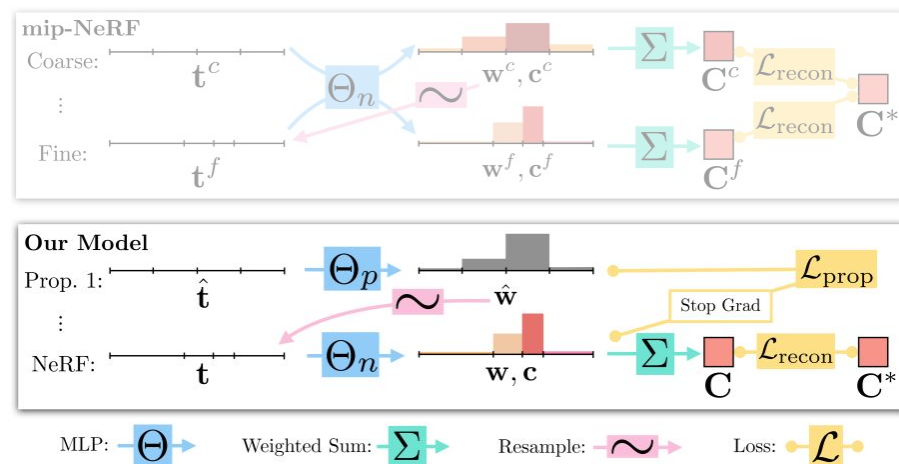
- NeRF: Representing scenes as neural radiance fields for view synthesis [Mildenhall et al., 2021]
- Mip-NeRF: A Multiscale Representation for Anti-Aliasing Neural Radiance Fields [Barron et al., 2021]



Sampling Application: Neural Rendering

Inverse transform sampling:

- NeRF: Representing scenes as neural radiance fields for view synthesis [Mildenhall et al., 2021]
- Mip-NeRF: A Multiscale Representation for Anti-Aliasing Neural Radiance Fields [Barron et al., 2021]



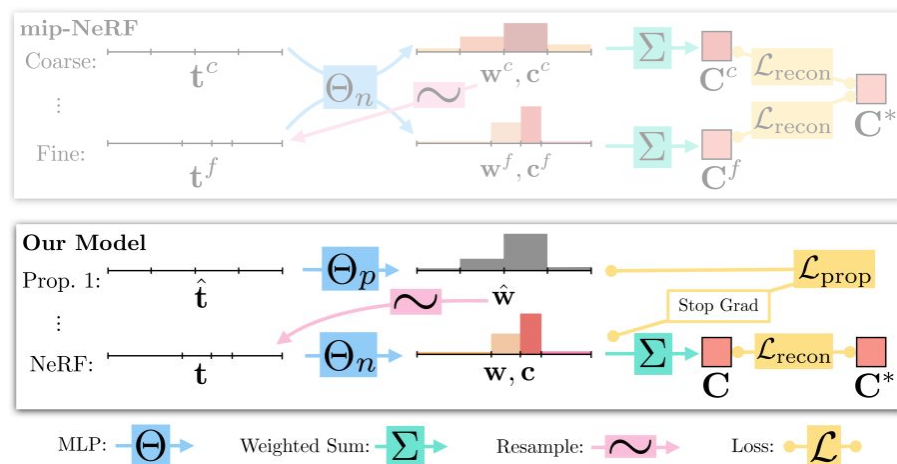
Proposal Network:

- Mip-NeRF 360: Unbounded Anti-Aliased Neural Radiance Fields [Barron et al., 2022]
- Zip-NeRF: Anti-aliased grid-based neural radiance fields [Barron et al., 2022]

Sampling Application: Neural Rendering

Inverse transform sampling:

- NeRF: Representing scenes as neural radiance fields for view synthesis [Mildenhall et al., 2021]
- Mip-NeRF: A Multiscale Representation for Anti-Aliasing Neural Radiance Fields [Barron et al., 2021]



Proposal Network:

- Mip-NeRF 360: Unbounded Anti-Aliased Neural Radiance Fields [Barron et al., 2022]
- Zip-NeRF: Anti-Aliased Grid-Based Neural Radiance Fields [Barron et al., 2022]



Particle Filter Derivation

Bayes' Rule

$$p(x|y, e) = \frac{p(y|x, e)p(x|e)}{p(y|e)}$$

$$bel(x_{0:t}) = p(x_{0:t} | u_{1:t}, z_{1:t})$$

$$p(x_{0:t} | z_{1:t}, u_{1:t})$$

$$\stackrel{\text{Bayes}}{=} \eta p(z_t | x_{0:t}, z_{1:t-1}, u_{1:t}) p(x_{0:t} | z_{1:t-1}, u_{1:t})$$

$$\stackrel{\text{Markov}}{=} \eta p(z_t | x_t) p(x_{0:t} | z_{1:t-1}, u_{1:t})$$

$$= \eta p(z_t | x_t) p(x_t | x_{0:t-1}, z_{1:t-1}, u_{1:t}) p(x_{0:t-1} | z_{1:t-1}, u_{1:t})$$

$$\stackrel{\text{Markov}}{=} \eta p(z_t | x_t) p(x_t | x_{t-1}, u_t) p(x_{0:t-1} | z_{1:t-1}, u_{1:t-1})$$

Particle Filter Derivation

Bayes' Rule

$$p(x|y, e) = \frac{p(y|x, e)p(x|e)}{p(y|e)}$$

$$bel(x_{0:t}) = p(x_{0:t} | u_{1:t}, z_{1:t})$$

$$p(x_{0:t} | z_{1:t}, u_{1:t})$$

$$\stackrel{\text{Bayes}}{=} \eta p(z_t | x_{0:t}, z_{1:t-1}, u_{1:t}) p(x_{0:t} | z_{1:t-1}, u_{1:t})$$

$$\stackrel{\text{Markov}}{=} \eta p(z_t | x_t) p(x_{0:t} | z_{1:t-1}, u_{1:t})$$

$$= \eta p(z_t | x_t) p(x_t | x_{0:t-1}, z_{1:t-1}, u_{1:t}) p(x_{0:t-1} | z_{1:t-1}, u_{1:t})$$

$$\stackrel{\text{Markov}}{=} \eta p(z_t | x_t) p(x_t | x_{t-1}, u_t) p(x_{0:t-1} | z_{1:t-1}, u_{1:t-1})$$

Particle Filter Derivation

Bayes' Rule

$$p(x|y, e) = \frac{p(y|x, e)p(x|e)}{p(y|e)}$$

$$bel(x_{0:t}) = p(x_{0:t} | u_{1:t}, z_{1:t})$$

$$p(x_{0:t} | z_{1:t}, u_{1:t})$$

$$\stackrel{\text{Bayes}}{=} \eta p(z_t | x_{0:t}, z_{1:t-1}, u_{1:t}) p(x_{0:t} | z_{1:t-1}, u_{1:t})$$

$$\stackrel{\text{Markov}}{=} \eta p(z_t | x_t) p(x_{0:t} | z_{1:t-1}, u_{1:t})$$

$$= \eta p(z_t | x_t) p(x_t | x_{0:t-1}, z_{1:t-1}, u_{1:t}) p(x_{0:t-1} | z_{1:t-1}, u_{1:t})$$

$$\stackrel{\text{Markov}}{=} \eta p(z_t | x_t) p(x_t | x_{t-1}, u_t) p(x_{0:t-1} | z_{1:t-1}, u_{1:t-1})$$

Particle Filter Derivation

Bayes' Rule

$$p(x|y, e) = \frac{p(y|x, e)p(x|e)}{p(y|e)}$$

$$bel(x_{0:t}) = p(x_{0:t} | u_{1:t}, z_{1:t})$$

$$p(x_{0:t} | z_{1:t}, u_{1:t})$$

$$\stackrel{\text{Bayes}}{=} \eta p(z_t | x_{0:t}, z_{1:t-1}, u_{1:t}) p(x_{0:t} | z_{1:t-1}, u_{1:t})$$

$$\stackrel{\text{Markov}}{=} \eta p(z_t | x_t) p(x_{0:t} | z_{1:t-1}, u_{1:t})$$

$$= \eta p(z_t | x_t) p(x_t | x_{0:t-1}, z_{1:t-1}, u_{1:t}) p(x_{0:t-1} | z_{1:t-1}, u_{1:t})$$

$$\stackrel{\text{Markov}}{=} \eta p(z_t | x_t) p(x_t | x_{t-1}, u_t) p(x_{0:t-1} | z_{1:t-1}, u_{1:t-1})$$

$$P(A, B) = P(A | B) P(B)$$

$$P(A, B | C) = P(A | B, C) P(B | C)$$

Particle Filter Derivation

Bayes' Rule

$$p(x|y, e) = \frac{p(y|x, e)p(x|e)}{p(y|e)}$$

$$bel(x_{0:t}) = p(x_{0:t} | u_{1:t}, z_{1:t})$$

$$p(x_{0:t} | z_{1:t}, u_{1:t})$$

$$\stackrel{\text{Bayes}}{=} \eta p(z_t | x_{0:t}, z_{1:t-1}, u_{1:t}) p(x_{0:t} | z_{1:t-1}, u_{1:t})$$

$$\stackrel{\text{Markov}}{=} \eta p(z_t | x_t) p(x_{0:t} | z_{1:t-1}, u_{1:t})$$

$$= \eta p(z_t | x_t) p(x_t | x_{0:t-1}, z_{1:t-1}, u_{1:t}) p(x_{0:t-1} | z_{1:t-1}, u_{1:t})$$

$$\stackrel{\text{Markov}}{=} \eta p(z_t | x_t) p(x_t | x_{t-1}, u_t) p(x_{0:t-1} | z_{1:t-1}, u_{1:t-1})$$

Particle Filter Derivation

Bayes' Rule

$$p(x|y, e) = \frac{p(y|x, e)p(x|e)}{p(y|e)}$$

$$bel(x_{0:t}) = p(x_{0:t} | u_{1:t}, z_{1:t})$$

$$p(x_{0:t} | z_{1:t}, u_{1:t})$$

$$\stackrel{\text{Bayes}}{=} \eta p(z_t | x_{0:t}, z_{1:t-1}, u_{1:t}) p(x_{0:t} | z_{1:t-1}, u_{1:t})$$

$$\stackrel{\text{Markov}}{=} \eta p(z_t | x_t) p(x_{0:t} | z_{1:t-1}, u_{1:t})$$

$$= \eta p(z_t | x_t) p(x_t | x_{0:t-1}, z_{1:t-1}, u_{1:t}) p(x_{0:t-1} | z_{1:t-1}, u_{1:t})$$

$$\stackrel{\text{Markov}}{=} \eta p(z_t | x_t) p(x_t | x_{t-1}, u_t) p(x_{0:t-1} | z_{1:t-1}, u_{1:t-1})$$



$$bel(x_{0:t-1}) \quad \text{Recursion!}$$

Particle Filter Derivation

$$p(x_{0:t} \mid z_{1:t}, u_{1:t}) = \eta p(z_t \mid x_t) p(x_t \mid x_{t-1}, u_t) p(x_{0:t-1} \mid z_{1:t-1}, u_{1:t-1})$$

- Recall importance weighting

$$w_t^{[m]} = \frac{\text{target distribution}}{\text{proposal distribution}}$$

Particle Filter Derivation

$$p(x_{0:t} \mid z_{1:t}, u_{1:t}) = \eta p(z_t \mid x_t) p(x_t \mid x_{t-1}, u_t) p(x_{0:t-1} \mid z_{1:t-1}, u_{1:t-1})$$

- Recall importance weighting

$$w_t^{[m]} = \frac{\text{target distribution}}{\text{proposal distribution}} \longleftarrow \text{bel}(x_{0:t})$$

Particle Filter Derivation

$$p(x_{0:t} \mid z_{1:t}, u_{1:t}) = \eta p(z_t \mid x_t) p(x_t \mid x_{t-1}, u_t) p(x_{0:t-1} \mid z_{1:t-1}, u_{1:t-1})$$

- Recall importance weighting

$$w_t^{[m]} = \frac{\text{target distribution}}{\text{proposal distribution}} \quad \begin{array}{l} \longleftarrow \text{bel}(x_{0:t}) \\ \longleftarrow p(x_t \mid x_{t-1}, u_t) \text{bel}(x_{0:t-1}) \end{array}$$

Particle Filter Derivation

$$p(x_{0:t} \mid z_{1:t}, u_{1:t}) = \eta p(z_t \mid x_t) p(x_t \mid x_{t-1}, u_t) p(x_{0:t-1} \mid z_{1:t-1}, u_{1:t-1})$$

- Recall importance weighting

$$w_t^{[m]} = \frac{\text{target distribution}}{\text{proposal distribution}}$$

We don't know this (don't have samples... yet)

← $bel(x_{0:t})$

← $p(x_t \mid x_{t-1}, u_t) bel(x_{0:t-1})$

We do know this (can generate samples)

Particle Filter Derivation

$$p(x_{0:t} \mid z_{1:t}, u_{1:t}) = \eta p(z_t \mid x_t) p(x_t \mid x_{t-1}, u_t) p(x_{0:t-1} \mid z_{1:t-1}, u_{1:t-1})$$

- Recall importance weighting

$$\begin{aligned} w_t^{[m]} &= \frac{\text{target distribution}}{\text{proposal distribution}} && \begin{array}{l} \longleftarrow bel(x_{0:t}) \\ \longleftarrow p(x_t \mid x_{t-1}, u_t) bel(x_{0:t-1}) \end{array} \\ &= \frac{\eta p(z_t \mid x_t) p(x_t \mid x_{t-1}, u_t) p(x_{0:t-1} \mid z_{1:t-1}, u_{1:t-1})}{p(x_t \mid x_{t-1}, u_t) p(x_{0:t-1} \mid z_{0:t-1}, u_{0:t-1})} \\ &= \eta p(z_t \mid x_t) \end{aligned}$$


But we know the ratio!


Particle Filter Algorithm


$$p(x_{0:t} \mid z_{1:t}, u_{1:t}) = \eta p(z_t \mid x_t) p(x_t \mid x_{t-1}, u_t) p(x_{0:t-1} \mid z_{1:t-1}, u_{1:t-1})$$

```
1:  Algorithm Particle_filter( $\mathcal{X}_{t-1}, u_t, z_t$ ):
2:       $\bar{\mathcal{X}}_t = \mathcal{X}_t = \emptyset$ 
3:      for  $m = 1$  to  $M$  do
4:          sample  $x_t^{[m]} \sim p(x_t \mid u_t, x_{t-1}^{[m]})$  ← motion model
5:           $w_t^{[m]} = p(z_t \mid x_t^{[m]})$  ← sensor model
6:           $\bar{\mathcal{X}}_t = \bar{\mathcal{X}}_t + \langle x_t^{[m]}, w_t^{[m]} \rangle$ 
7:      endfor
8:      for  $m = 1$  to  $M$  do
9:          draw  $i$  with probability  $\propto w_t^{[i]}$  ← importance sampling
10:         add  $x_t^{[i]}$  to  $\mathcal{X}_t$ 
11:      endfor
12:      return  $\mathcal{X}_t$ 
```


\mathcal{X}_{t-1} - previous particle set
 \mathcal{X}_t - output particle set








Measurement z
 $p(z_t \mid x_t)$



Measurement z


$$p(z_t | x_t)$$






Measurement z

$$p(z_t | x_t)$$






Measurement z


$$p(z_t | x_t)$$






Measurement z

$$p(z_t | x_t)$$







Motion u

$$p(x_t \mid x_{t-1}, u_t)$$

Measurement z


$$p(z_t \mid x_t)$$






Motion u


$$p(x_t \mid x_{t-1}, u_t)$$



Measurement z


$$p(z_t \mid x_t)$$






Motion u


$$p(x_t \mid x_{t-1}, u_t)$$




Measurement z

$$p(z_t \mid x_t)$$





Motion u



$$p(x_t | x_{t-1}, u_t)$$


Measurement z




$$p(z_t | x_t)$$

What should the motion model look like?


- Needs to be a distribution from which we can draw samples
- Usually use a normal distribution with odometry measurement as mean




Sampling from the motion model



Sampling from the motion model




Sampling from the motion model




1D Example

- State uniformly distributed initially





1D Example

- Sensor model, update weights




1D Example

- Motion model, resample particles




1D Example

- Sensor model, update weights




2D Example

- Initially particles uniformly distributed




2D Example

- Sensor model, update weights, resample




2D Example

- Propagate with motion model




2D Example

- Get measurements




2D Example

- Get measurements




2D Example

- Resample




2D Example

- Propagate with motion model




2D Example

- Get measurements




2D Example

- Get measurements




Resample

- Resample




2D example

- Propagate with motion model



2D example


- Get measurements




and so forth...

What should the sensor model look like?

- Depends on the sensor
- Previous example uses a laser range finder (LRF), which sends out lasers at known angles and measures the range
- Therefore we have a 1-dimensional measurement which is conditioned on the estimated state and the map



Range Sensor Model




z_t^*
 z_t



$$p(z_t^k \mid x_t, m) ??$$

Range Sensor Model

- Gives likelihood of seeing measurement z_t , given:
 - State x_t
 - Map m




Range Sensor Model


Superposition of:

- Gaussian centered at predicted range according to map + particle
- Uniform at end-of-range
- Exponential decay
- Maybe another uniform over entire range
- ...


(a) Gaussian distribution p_{hit}




(b) Exponential distribution p_{short}



(c) Uniform distribution p_{max}



(d) Uniform distribution p_{rand}




Range Sensor Model




Superposition of:

- Gaussian centered at predicted range according to map + particle
- Uniform at end-of-range
- Exponential decay
- Maybe another uniform over entire range
- ...

(a) Sonar data



(b) Laser data



Particle Filter Algorithm

$$p(x_{0:t} \mid z_{1:t}, u_{1:t}) = \eta p(z_t \mid x_t) p(x_t \mid x_{t-1}, u_t) p(x_{0:t-1} \mid z_{1:t-1}, u_{1:t-1})$$

```
1:  Algorithm Particle_filter( $\mathcal{X}_{t-1}, u_t, z_t$ ):
2:       $\bar{\mathcal{X}}_t = \mathcal{X}_t = \emptyset$ 
3:      for  $m = 1$  to  $M$  do
4:          sample  $x_t^{[m]} \sim p(x_t \mid u_t, x_{t-1}^{[m]})$  ← motion model
5:           $w_t^{[m]} = p(z_t \mid x_t^{[m]})$  ← sensor model
6:           $\bar{\mathcal{X}}_t = \bar{\mathcal{X}}_t + \langle x_t^{[m]}, w_t^{[m]} \rangle$ 
7:      endfor
8:      for  $m = 1$  to  $M$  do
9:          draw  $i$  with probability  $\propto w_t^{[i]}$  ← importance sampling
10:         add  $x_t^{[i]}$  to  $\mathcal{X}_t$ 
11:      endfor
12:      return  $\mathcal{X}_t$ 
```

\mathcal{X}_{t-1} - previous particle set
 \mathcal{X}_t - output particle set

Implementation problems

- Stationary robot
 - No motion or sensor data
 - Converges to one particle (likely incorrect)
- Particle deprivation
 - Too few particles
 - “Correct” particles die out
- Very precise sensor models can result in poor performance
 - Accurate odometry, poor range sensor?
 - Poor odometry, accurate range sensor?

Tips & Tricks


- Low variance resampling



- Store logs of weights
- Experiment with different sensor models

Tips & Tricks

- Low variance resampling



- Store logs of weights
- Experiment with different sensor models

Tips & Tricks

- Low variance resampling





Figure 4.3 Principle of the low variance resampling procedure. We choose a random number r and then select those particles that correspond to $u = r + (m - 1) \cdot M^{-1}$ where $m = 1, \dots, M$.

- Store logs of weights
- Experiment with different sensor models

FastSLAM

- FastSLAM: A Factored Solution to the Simultaneous Localization and Mapping Problem [Montemerlo et al., 2002]



Localization vs. SLAM

- A particle filter can be used to solve both problems
- Localization: state space $\langle x, y, \theta \rangle$
- SLAM: state space $\langle x, y, \theta, \text{map} \rangle$
 - for landmark maps = $\langle l_1, l_2, \dots, l_m \rangle$
 - for grid maps = $\langle c_{11}, c_{12}, \dots, c_{1n}, c_{21}, \dots, c_{nm} \rangle$
- **Problem:** The number of particles needed to represent a posterior grows exponentially with the dimension of the state space!

Dependencies

- Localization:

$$p(x_{0:t} \mid z_{1:t}, u_{1:t}) = \eta p(z_t \mid x_t) p(x_t \mid x_{t-1}, u_t) p(x_{0:t-1} \mid z_{1:t-1}, u_{1:t-1})$$

- Is there a dependency between the dimensions of the state space?
- If so, can we use the dependency to solve the problem more efficiently?
- In the SLAM context
 - The map depends on the poses of the robot.
 - We know how to build a map given the position of the sensor is known.

Factored Posterior (Landmarks)

Diagram illustrating the factored posterior for SLAM with landmarks. The equation is:

$$p(x_{1:t}, l_{1:m} \mid z_{1:t}, u_{0:t-1}) = p(x_{1:t} \mid z_{1:t}, u_{0:t-1}) \cdot p(l_{1:m} \mid x_{1:t}, z_{1:t})$$

Labels and arrows indicating dependencies:

- poses** (blue text) points down to the first term $p(x_{1:t} \mid z_{1:t}, u_{0:t-1})$.
- map** (blue text) points down to the second term $p(l_{1:m} \mid x_{1:t}, z_{1:t})$.
- observations & control inputs** (blue text) points down to the first term $p(x_{1:t} \mid z_{1:t}, u_{0:t-1})$ and diagonally down to the second term $p(l_{1:m} \mid x_{1:t}, z_{1:t})$.
- SLAM posterior** (black text) points up to the entire equation.
- Robot path posterior** (black text) points up to the first term $p(x_{1:t} \mid z_{1:t}, u_{0:t-1})$.
- landmark positions** (black text) points up to the second term $p(l_{1:m} \mid x_{1:t}, z_{1:t})$.

Does this help to solve the problem?


Factored Posterior (Landmarks)

The diagram illustrates the factorization of the SLAM posterior. At the top, three labels in blue text—'poses', 'map', and 'observations & control inputs'—have blue arrows pointing down to the first three arguments of the SLAM posterior equation. The equation is $p(x_{1:t}, l_{1:m} \mid z_{1:t}, u_{0:t-1}) = p(x_{1:t} \mid z_{1:t}, u_{0:t-1}) \cdot p(l_{1:m} \mid x_{1:t}, z_{1:t})$. Below the equation, three labels have arrows pointing up to their respective parts: 'SLAM posterior' points to the left-hand side, 'Robot path posterior' points to $p(x_{1:t} \mid z_{1:t}, u_{0:t-1})$, and 'landmark positions' points to $p(l_{1:m} \mid x_{1:t}, z_{1:t})$.

$$\begin{array}{c} \text{poses} \quad \text{map} \quad \text{observations \& control inputs} \\ \downarrow \quad \downarrow \quad \downarrow \quad \swarrow \\ p(x_{1:t}, l_{1:m} \mid z_{1:t}, u_{0:t-1}) = \\ \uparrow \quad \uparrow \quad \uparrow \\ \text{SLAM posterior} \quad p(x_{1:t} \mid z_{1:t}, u_{0:t-1}) \cdot p(l_{1:m} \mid x_{1:t}, z_{1:t}) \\ \text{Robot path posterior} \quad \text{landmark positions} \end{array}$$

Solve localization and mapping separately!


Mapping using Landmarks




Knowledge of the robot's true path renders landmark positions conditionally independent

Factored Posterior

$$\begin{aligned} & p(x_{1:t}, l_{1:m} \mid z_{1:t}, u_{0:t-1}) \\ &= p(x_{1:t} \mid z_{1:t}, u_{0:t-1}) \cdot p(l_{1:m} \mid x_{1:t}, z_{1:t}) \\ &= p(x_{1:t} \mid z_{1:t}, u_{0:t-1}) \cdot \prod_{i=1}^M p(l_i \mid x_{1:t}, z_{1:t}) \end{aligned}$$



Robot path posterior
(localization problem)




Conditionally independent
landmark positions

Rao-Blackwellization


$$p(x_{1:t}, l_{1:m} \mid z_{1:t}, u_{0:t-1}) = \\ p(x_{1:t} \mid z_{1:t}, u_{0:t-1}) \cdot \prod_{i=1}^M p(l_i \mid x_{1:t}, z_{1:t})$$

- This factorization is also called Rao-Blackwellization
- Given that the second term can be computed efficiently, particle filtering becomes possible!




FastSLAM

- Rao-Blackwellized particle filtering based on landmarks [Montemerlo et al., 2002]
- Each landmark is represented by a 2x2 Extended Kalman Filter (EKF)
- Each particle therefore has to maintain M EKFs




FastSLAM - Action Update




50

FastSLAM - Sensor Update



51

FastSLAM - Sensor Update



52

FastSLAM Complexity

- Update robot particles based on control u_{t-1}
- Incorporate observation z_t into Kalman filters
- Resample particle set

N = Number of particles

M = Number of map features

$$O(N)$$

Constant time per particle

$$O(N \cdot \log(M))$$

log time per particle

$$O(N \cdot \log(M))$$

log time per particle

$$O(N \cdot \log(M))$$

log time per particle

FastSLAM Complexity

- Update robot particles based on control u_{t-1}
- Incorporate observation z_t into Kalman filters
- Resample particle set

N = Number of particles

M = Number of map features

$$O(N)$$

Constant time per particle

$$O(N \cdot \log(M))$$

log time per particle

$$O(N \cdot \log(M))$$

log time per particle

$$O(N \cdot \log(M))$$

log time per particle

$$\text{EKF SLAM: } O(K^{2.8} + N^2)$$

K = Measurement vector dimension


N = State vector dimension

Results - Victoria Park

- 4 km traverse
- < 5 m RMS position error
- 100 particles


Blue = GPS

Yellow = FastSLAM




Dataset courtesy of University of Sydney

Results - Number of Particles



Results - Motion Uncertainty



FastSLAM slides courtesy of

Sebastian Thrun, Wolfram Burgard, Dieter
Fox

Publicly available at
www.probablistic-robotics.org